




Java course - IAG0040

Threads & Concurrency

Introduction to Threads

- A *Thread* is a lightweight process
- A *Thread* is a single sequential flow of control within a program



More formal definition

- A **thread** is pure activity, i.e. the information about the current position of execution (instruction pointer), the registers and the current stack
- An **address space** is a range of addressable memory that is only accessible by activities that are allowed to read/write from it
- A **process** consists of an address space and one or multiple threads

Threads & Java

- Java supports Threads and synchronization natively
- Threads can be used for doing two or more tasks at once
- `main()` method is always executed in the “main” thread
 - there are always several more system threads running in your Java program, e.g. garbage collector
- You can always start as many new threads as your program needs using the *Thread* class

Thread class

- Threads run *Runnable* tasks, there are two possible implementations
 - extend the *Thread* class itself, overriding the run method:

```
class MegaThread extends Thread { ... }  
Thread t = new MegaThread();
```
 - implement the *Runnable* interface and pass it to a new *Thread*:

```
class MegaRunnable implements Runnable { ... }  
Thread t = new Thread(new MegaRunnable());
```
- Every *Thread* has a name: it can be specified on creation or using the `setName()` method
- *Thread* class provides some useful static methods:
 - `currentThread()` returns the current Thread's instance
 - `sleep()` pauses the current Thread for the given time
 - `yield()` allows other Threads to run by pausing the current one

Thread class (cont)

- `start()` starts the new Thread in the background, note: `run()` method doesn't start the Thread!
- `join()` waits for another Thread to die
- `setDaemon(true)` sets the daemon flag. JVM terminates if there are only daemon threads left running.
- `getId()` returns a unique Id (long) of the Thread
- `isAlive()` tests if the Thread is alive (started and not dead)
- `getState()` provides more information on the state
- There are many other interesting `is/get` methods, see Javadoc

Exception handling

- Each thread has a default top-level exception handler
 - that calls `printStackTrace()`
- Just like the **main** thread, where the `main()` method is executed
- If an exception is thrown out of a thread, the thread is **terminated**

Interrupting Threads

- There are many deprecated methods in the Thread class, which have been proved to be dangerous (may cause deadlocks, etc): `suspend()`, `resume()`, `stop()`, `destroy()`, etc
- Instead, `interrupt()` can be used
 - it interrupts blocking/waiting method calls, e.g. `sleep()`, `wait()`, `join()`, etc, with an `InterruptedException`
 - it signals the Thread that it should terminate
 - interrupted flag can be checked using either the static `Thread.interrupted()` or `isInterrupted()` method
 - it is a safe and graceful way to interrupt a Thread, because it involves interaction from the Thread itself

Scheduling

- Scheduling is execution of multiple threads on a single CPU, dividing available CPU time into **time slices**
- Multiple threads are the only possibility for a program to use multiple CPUs / cores (if they are available)
- Each Thread has a priority between `MIN_PRIORITY` and `MAX_PRIORITY` (currently from 1 to 10)
- Java implements **fixed-priority preemptive scheduling**
 - at any given time, a Thread with the highest priority is running, but this is not guaranteed: lower-priority threads may be executed to avoid starvation
 - if any higher-priority Thread appears, it is executed pausing others

Synchronization

- Independent and asynchronous Threads are fine without synchronization
- However, access to shared data must be synchronized
 - thread scheduling is unpredictable: threads may access (read-modify-write) variables in any order
- Synchronization prevents interruption of critical regions in the code
- Dijkstra's *Dining Philosophers Problem* is often used for illustration of various synchronization/concurrency problems

Synchronization (cont)

- Java provides the **synchronized** keyword
 - it is used to obtain a lock on an arbitrary object instance
 - you can declare **synchronized** blocks of code:
`synchronized (lockObject) { ... }`
 - code within the **synchronized** block is accessed atomically
 - you can also declare methods as **synchronized** - the lock is obtained on the instance of the object or the Class object if method is **static**

Synchronization (cont)

- Object class provides `wait()` and `notify()` methods
 - both require acquiring a lock on the object first (**synchronized**)
 - they can be used for signaling to different parts in the code
 - `wait()` can have a timeout and may be interrupted
- Java 1.5 provides many exciting new features for concurrent programming in the **java.util.concurrent** package
 - *Lock and ReadWriteLock* implementations (e.g. *ReentrantLock*), *Semaphore*, etc provide additional features and sometimes better scalability
 - ```
ReentrantLock rlock = new ReentrantLock();
rlock.lock();
try { ... } finally { rlock.unlock(); }
```



# Thread safety

- Rule #1: Document thread safety in Javadoc!!!
- Many classes or methods can have various levels of thread safety
  - **Immutable** - always thread-safe, because its state never changes
  - **Thread-safe** - can be used in multiple threads concurrently (either synchronization is not needed or it is synchronized internally), this is a rather strict requirement
  - **Conditionally thread-safe** - each individual operation may be thread safe, but certain sequences of operations may require external synchronization, e.g. *Iterator* returned by *Vector* and *Hashtable*
  - **Thread compatible** - not thread safe, but thread safety can be achieved by external synchronization of method calls, e.g. *StringBuilder*, all *Collections*, etc
  - **Thread hostile** - external synchronization of individual method calls will not make an object thread safe, this is rare and can be a result of bad design



# Timers and Tasks

- *Timer* and *TimerTask* (both in `java.util`) can be used for conveniently scheduling task execution at intervals or with delays
- Extend the abstract *TimerTask* to create your task (implement the `run()` method)
- Use a *Timer* for scheduling this task arbitrarily. Every instance of a *Timer* represents a single Thread.
  - `schedule()` - schedules a task for either single or periodical execution. Counting to the next execution begins after the previous one is finished.
  - `scheduleAtFixedRate()` - schedules with the exact period. Counting to the next execution is not affected by the previous one.
  - any *Timer* can have many *TimerTasks* scheduled



# ThreadLocal

- *java.lang.ThreadLocal* can be used for storing independent data for each Thread (thread-local variables)
  - All Threads share the same *ThreadLocal* instance, but every *Thread* only 'sees' data belonging to it
  - Can be used for storing such things, as User ID, Transaction ID, etc which are handled by current *Thread*
- *ThreadLocal* methods
  - `set(...)` - sets the thread-local variable
  - `get()` - retrieves the previously set thread-local variable



# More on java.util.concurrent

- Introduced in Java 1.5, besides higher-level synchronization classes it provides many useful functionality for reduced programming effort, more performance, reliability, maintainability, and productivity
- Task scheduling framework: the *Executor* framework standardizes scheduling, invocation, execution, and control of asynchronous *Runnable* tasks (implementation includes thread pools)
- Concurrent Collections: some new implementations of Map, List, and Queue
- Atomic variables: **atomic** subpackage provides classes for atomically manipulating of variables, e.g. *AtomicInteger* (for higher performance)
- Synchronizers and Locks: more general purpose implementations with timeouts, reader/writer differentiation, non-nested scoping, etc
- Nanosecond-granularity timing, e.g. `System.nanoTime()` ;